Forecasting S & P 500 Using SARIMAX

Matthew Chen, William Gipson,Ann Kim, Liangyong Wang

Goals & Objectives

Project Goals and Objectives

Goals:

- Forecast daily S & P 500 prices using SARIMAX with exogenous variables

Objectives:

- Incorporate relevant exogenous variables
- Apply proper stationarity transformation and dimensionality reduction
- Use rolling-origin backtesting to evaluate model accuracy

Background

What is the S & P 500?

- List of 500 largest U.S. companies
- Used to measure stock market health
- Helps track U.S. economic trends
- Includes Apple, Microsoft, and others
- Reacts to global news and events

SARIMAX: Seasonal ARIMA with Exogenous Variables

Model Structure: SARIMA (p, d, q) x (P, D, Q)s + X

- **AR (p)**: Autoregressive terms
- **I (d)**: Differencing
- **MA (q)**: Moving average terms
- **Seasonal (P, D, Q)**: Seasonal AR, differencing, and MA terms with period s
- **X**: Exogenous variables

When to Use SARIMAX:

- **trend and/or seasonality are detected**
- **External factors are detected**

Common SARIMAX Applications:

- **Forecasting retail sales**
- **Modeling electricity demand**

Why Use SARIMAX for Forecasting the S&P 500?

- Forecasting its movement is important for investors and policymakers.
- Traditional ARIMA models do not account for external factors.
- SARIMAX allows the use of exogenous variables.

Examples Used in Our Project

Summary of Reference Paper (Erlemann et al., 2025)

Title: SARIMAX-Based Framework for S&P 500 Forecasting: Incorporating Economics Indicators

Published: May 2025 (Preprint on ResearchGate)

- S&P 500 is inherently non-stationary
- Acknowledge technological impact to the economy

Why We Use This Paper

Relevance to Our Project:

- Same target: S&P 500 daily forecasting
- Same model: SARIMAX with external variables

Benefits

- Clear, tested pipeline to follow
- Benchmarks we can compare against
- Real-world exogenous features + strong validation method

How Our Work Goes Beyond the Paper

- ❏ Parallel Processing
- ❏ Residual diagnostic
- ❏ Kalman smoothing
- ❏ Mean Absolute Percentage Error (MAPE)
- ❏ Forecast visualization

Data Components & Modeling

Exogenous Variables Used in SARIMAX Model

Category	Exogenous Variable Included
Financial Indicators	Gold, Crude Oil, Copper, Bitcoin, VIX, S&P GSCI
Treasury Bond & Currency Exchange	Treasury Yields (3M, 10Yr), forex(foreign currency exchange): EUR — USD, GBP—USD, JPY —USD, AUD—USD, CAD—USD
Stock Indices	SSE(Shanghai), STOXX 600 (Europe), MOEX (Moscow)
Google Trends	“SP500”, “ETF”, “Index Fund”, “SPX”
Unemployment	United States Unemployment Rate

Target Variable Used in SARIMAX Model

- Target Variable: daily closing value of the S&P 500 index
- Compare Predicted vs. Actual time series models of the target variable
- GOAL: small differences between the predicted and actual time series model
- Accuracy Analysis Measure: R^2, MAE, RMSE, MAPE

SARIMAX Pipeline

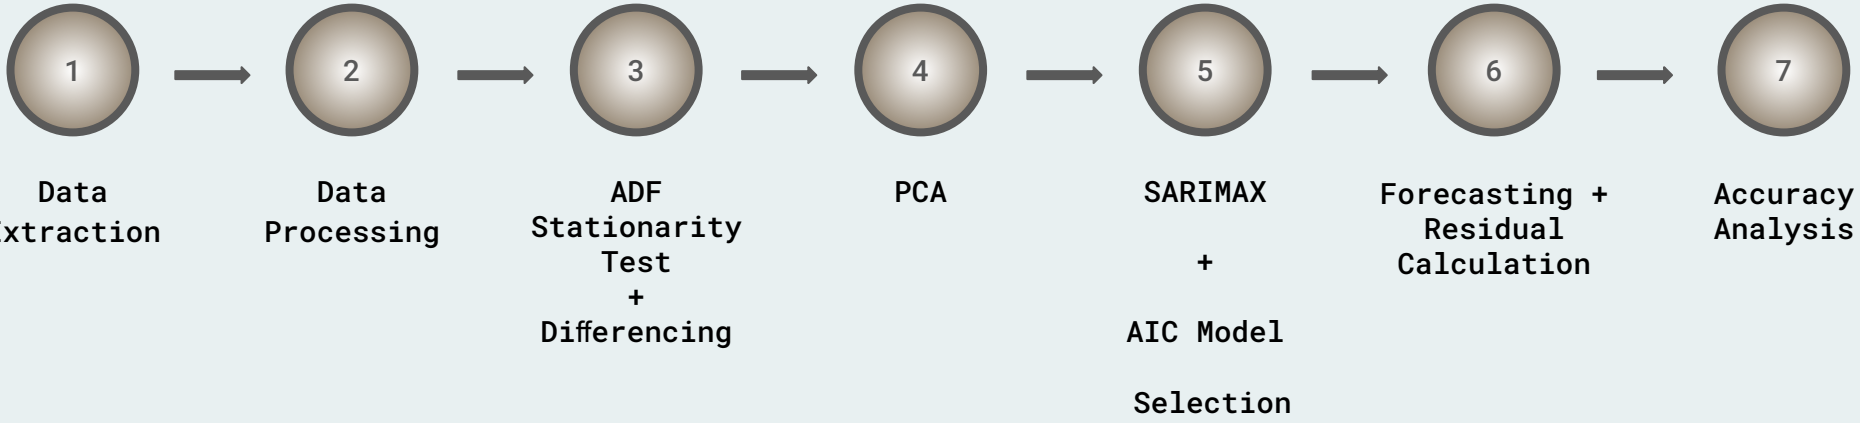

Handling Missing data & Imputation

Why it mattered

- Multiple macro-financial series came from different sources → uneven date coverage & sporadic gaps
- Stationary models (ADF + SARIMAX) need complete, numeric inputs

Imputation strategy

- **Primary fill**:
  - Kalman smoother built on an automatically-selected ARIMA for each series
  - Reconstructs values consistent with each series' own dynamics
- **Edge repair**:
  - **Forward LOCF** – pushes last known value forward
  - **Backward LOCF** – back-fills at the very start
  - Guarantees *no leading or trailing NA* in any regressor

Findings from Exploratory Data Analysis

- **Preprocessing**:
  - Merged all datasets by date
  - Removed all-NA or constant columns
  - Imputed missing values using Kalman smoothing and LOCF
- **Stationarity Check**:
  - Applied ADF tests on all series
  - Differenced non-stationary training time series
- **Dimensionality Reduction**:
  - Used PCA to reduce the number of exogenous features while retaining 95% of the variance
- **Initial Observations**:
  - Target variable shows an increasing general trend

Results

S & P 500: Actual vs SARIMAX

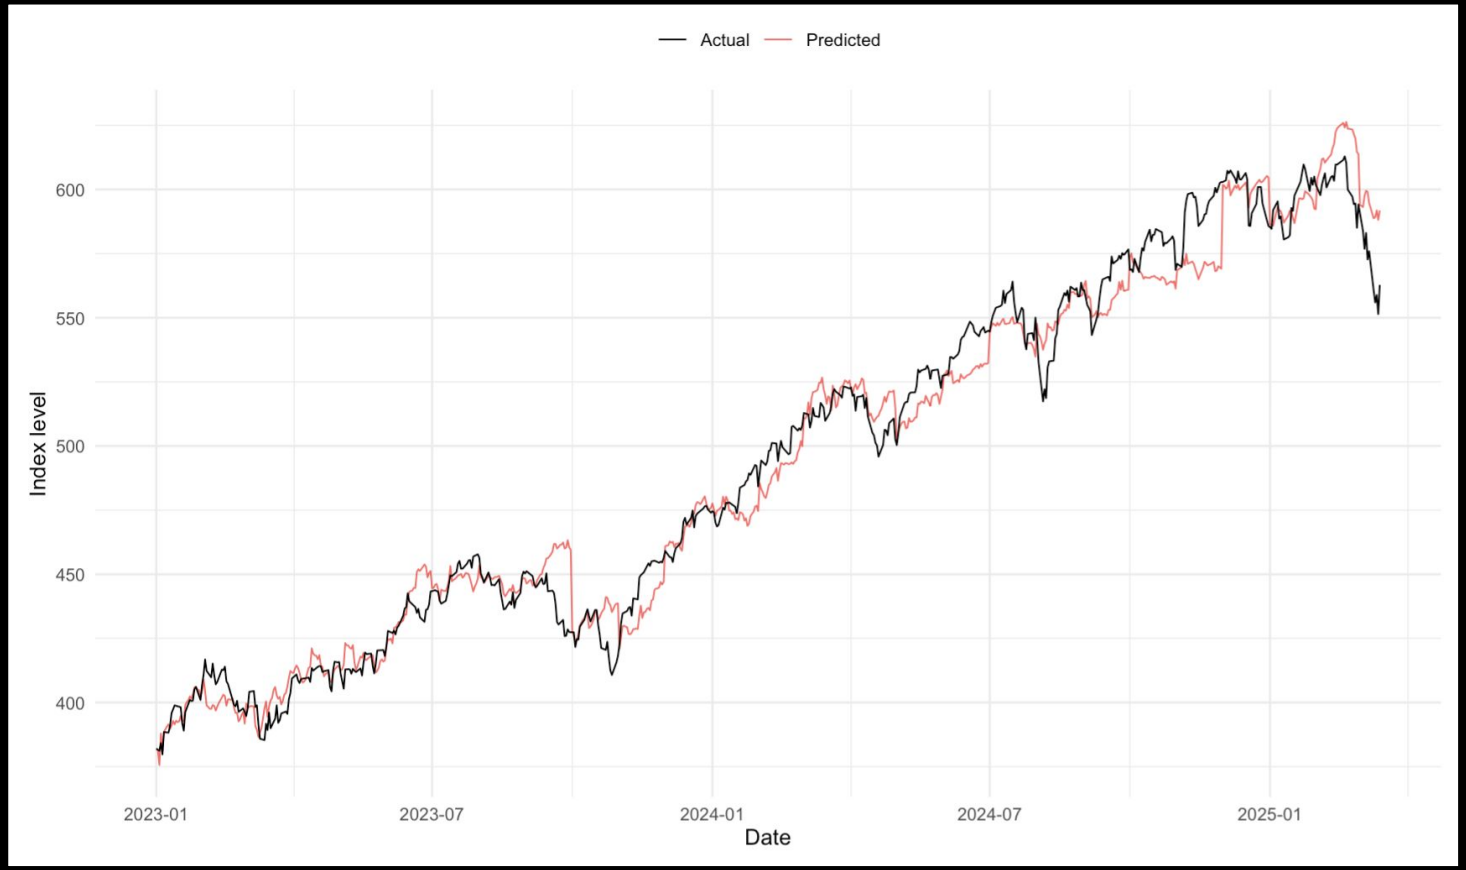

SARIMAX Model Residual Diagnostic Checks: Residuals Over Time

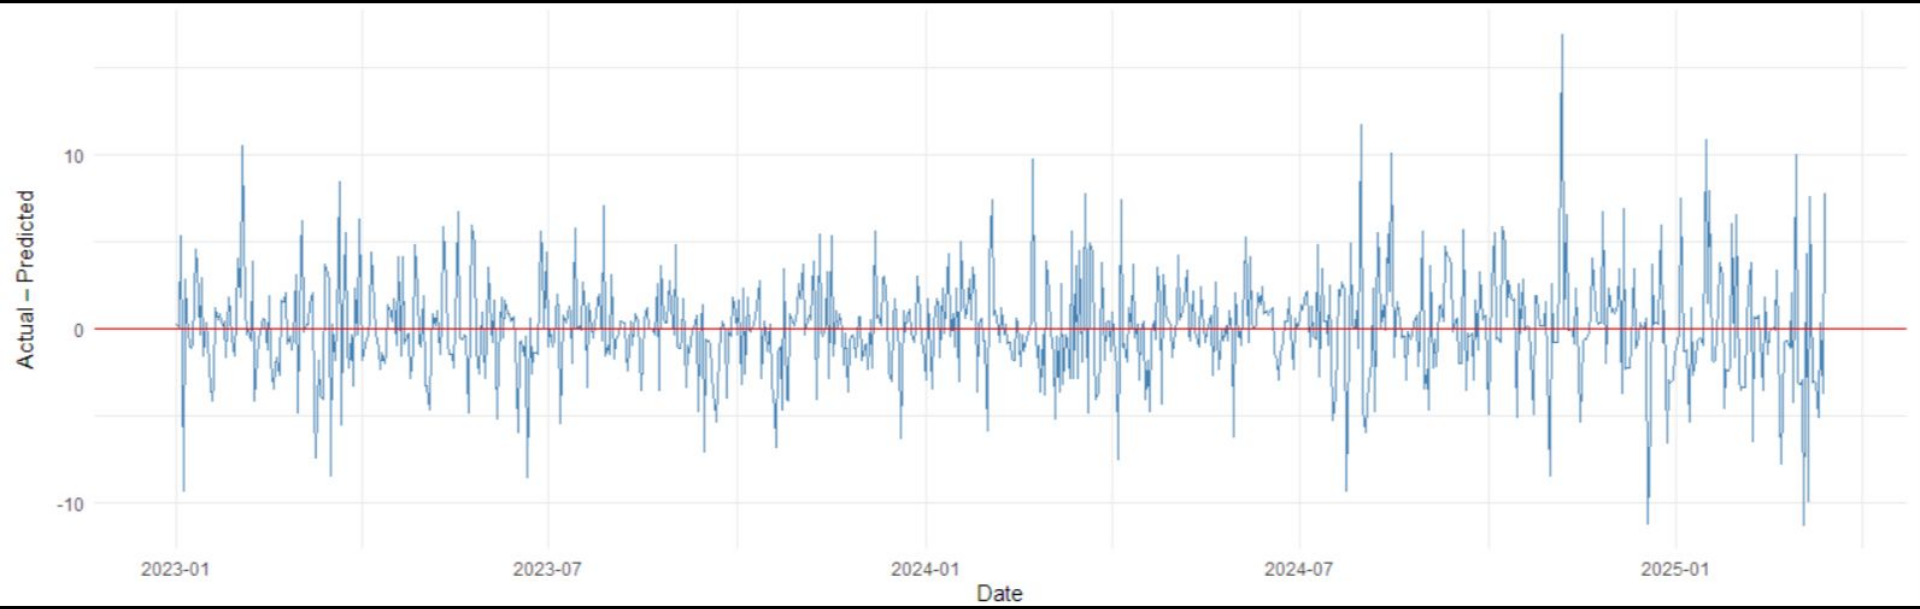

SARIMAX Model Residual Diagnostic Checks: Residual Distribution

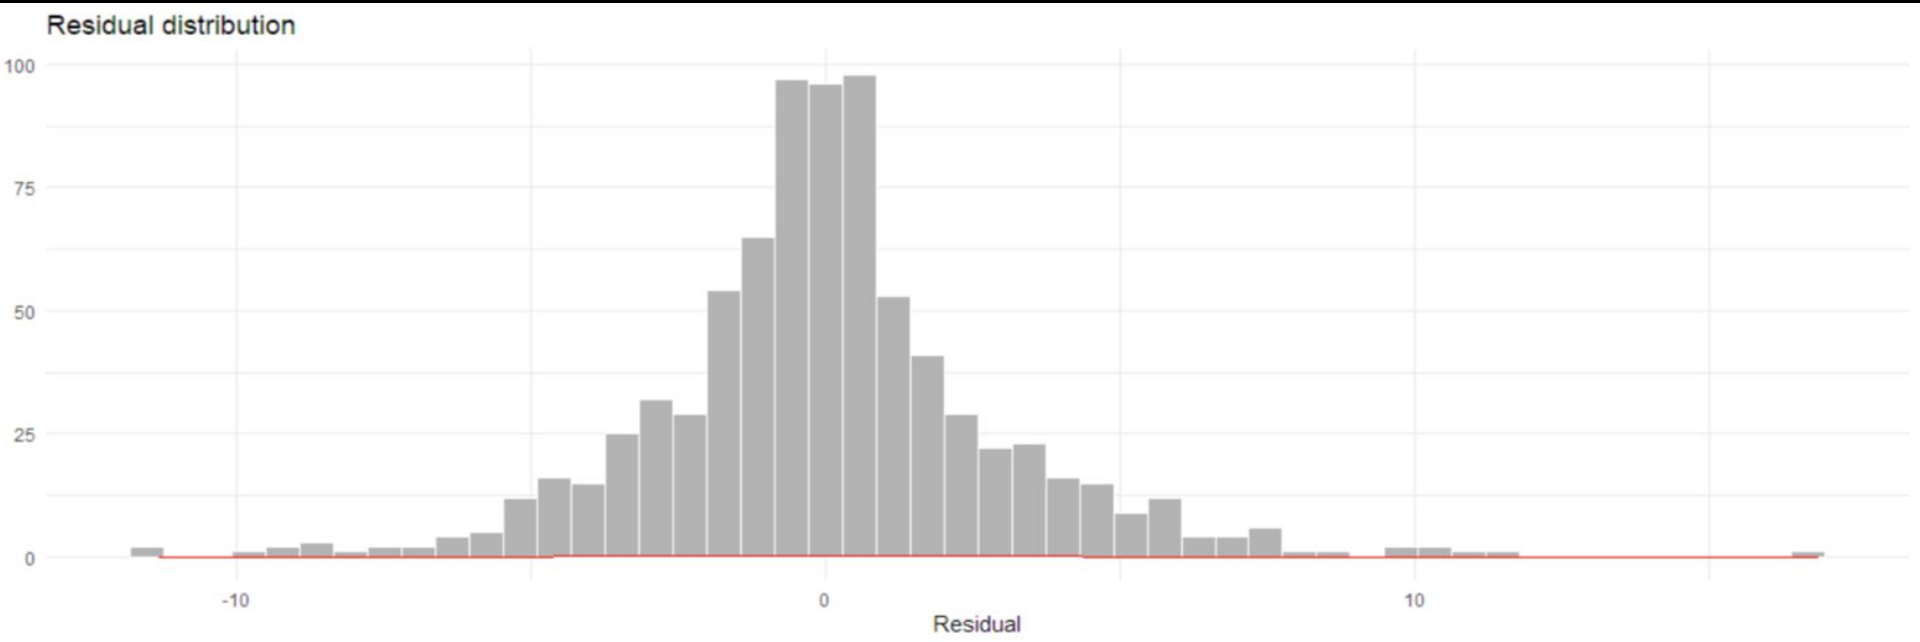

SARIMAX Model Residual Diagnostic Checks: ACF Plot

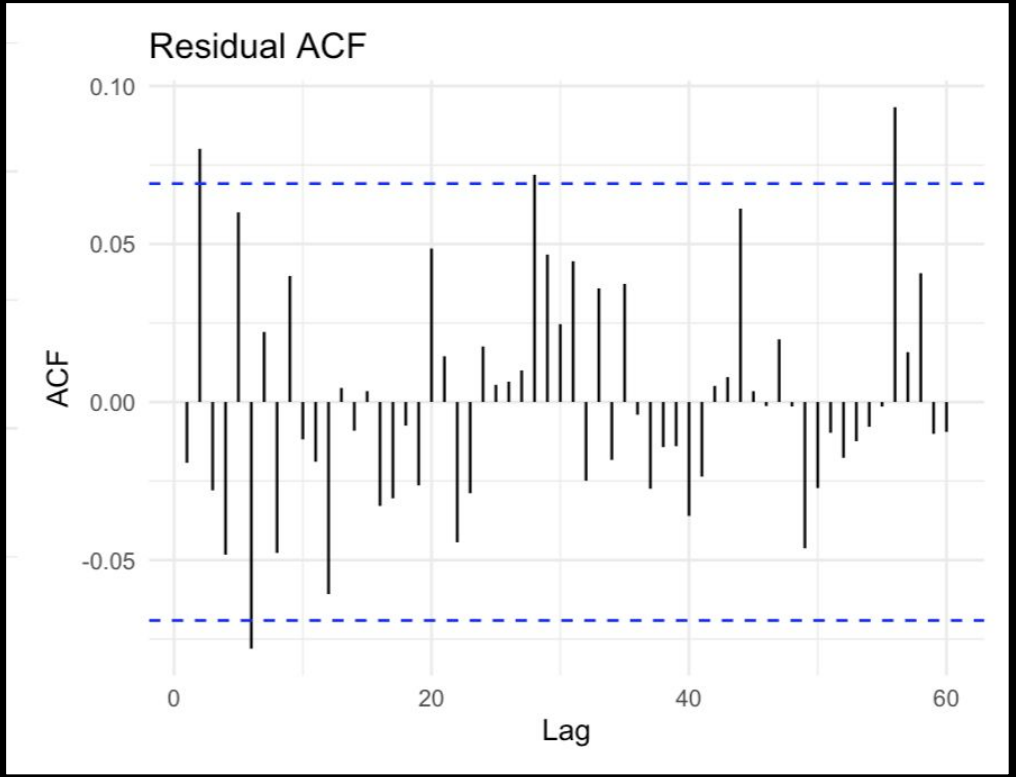

SARIMAX Performance Summary

- **What is the S&P 500 Index?**
  - A score that shows how 500 big U.S. companies are doing overall
  - Changes every second during market hours (Mon–Fri, 9:30am–4:00pm EST)
  - Today's values are around 4,000–5,500
- **MAE (Mean Absolute Error):**
  - Shows average mistake size in points
  - MAE = 8.375 → Forecast is off by about 8 index points on average
- **RMSE (Root Mean Squared Error):**
  - Like MAE but gives more weight to bigger mistakes
  - RMSE = 11.084 → Slightly higher average error when larger mistakes are penalized more
- **R^2 (R-squared)**
  - Proportion of variance explained by the model
  - R^2 = 0.974 → Model explains 97.4% of the S&P 500's variation
- **MAPE (Mean Absolute Percentage Error)**
  - Average percentage error
  - MAPE = 1.67% → Forecasts are off by just 1.67%, indicating excellent accuracy

Future Directions

- Incorporate tariff data into the SARIMAX model to improve performance (e.g., average percentage of tariffs)

- Compare the SARIMAX model to auto.arima() in R and apply variance stabilization transformations using the Box-Cox transformation parameter if necessary

Questions?